50.051 Programming Language Concepts

W11-S2 Bottom-Up Parsing (Part 1)

Matthieu De Mari

Different types of Parsing

Two big families of parsing algorithms

Top-Down Parsing

- Seen before, best we could do was in the case of LL(1) grammars, but very few grammars are going to be LL(k).

Bottom-Up Parsing

- Start from the input string *x*, whose syntax neds to be verified.

- Work you way back to a start symbol.

- Basically, the Top-Down parsing task, but in reverse!

- Builds on ideas of Top-Down parsing, but more efficient, and will work on non-LL(k) grammars.

Introduction example for Bottom-Up Parsing

Fact #1 (to be confirmed later): Bottom-Up Parsers can deal with non-LL(k) grammars and can also handle left-recursive grammars.

- Consider the following grammar:
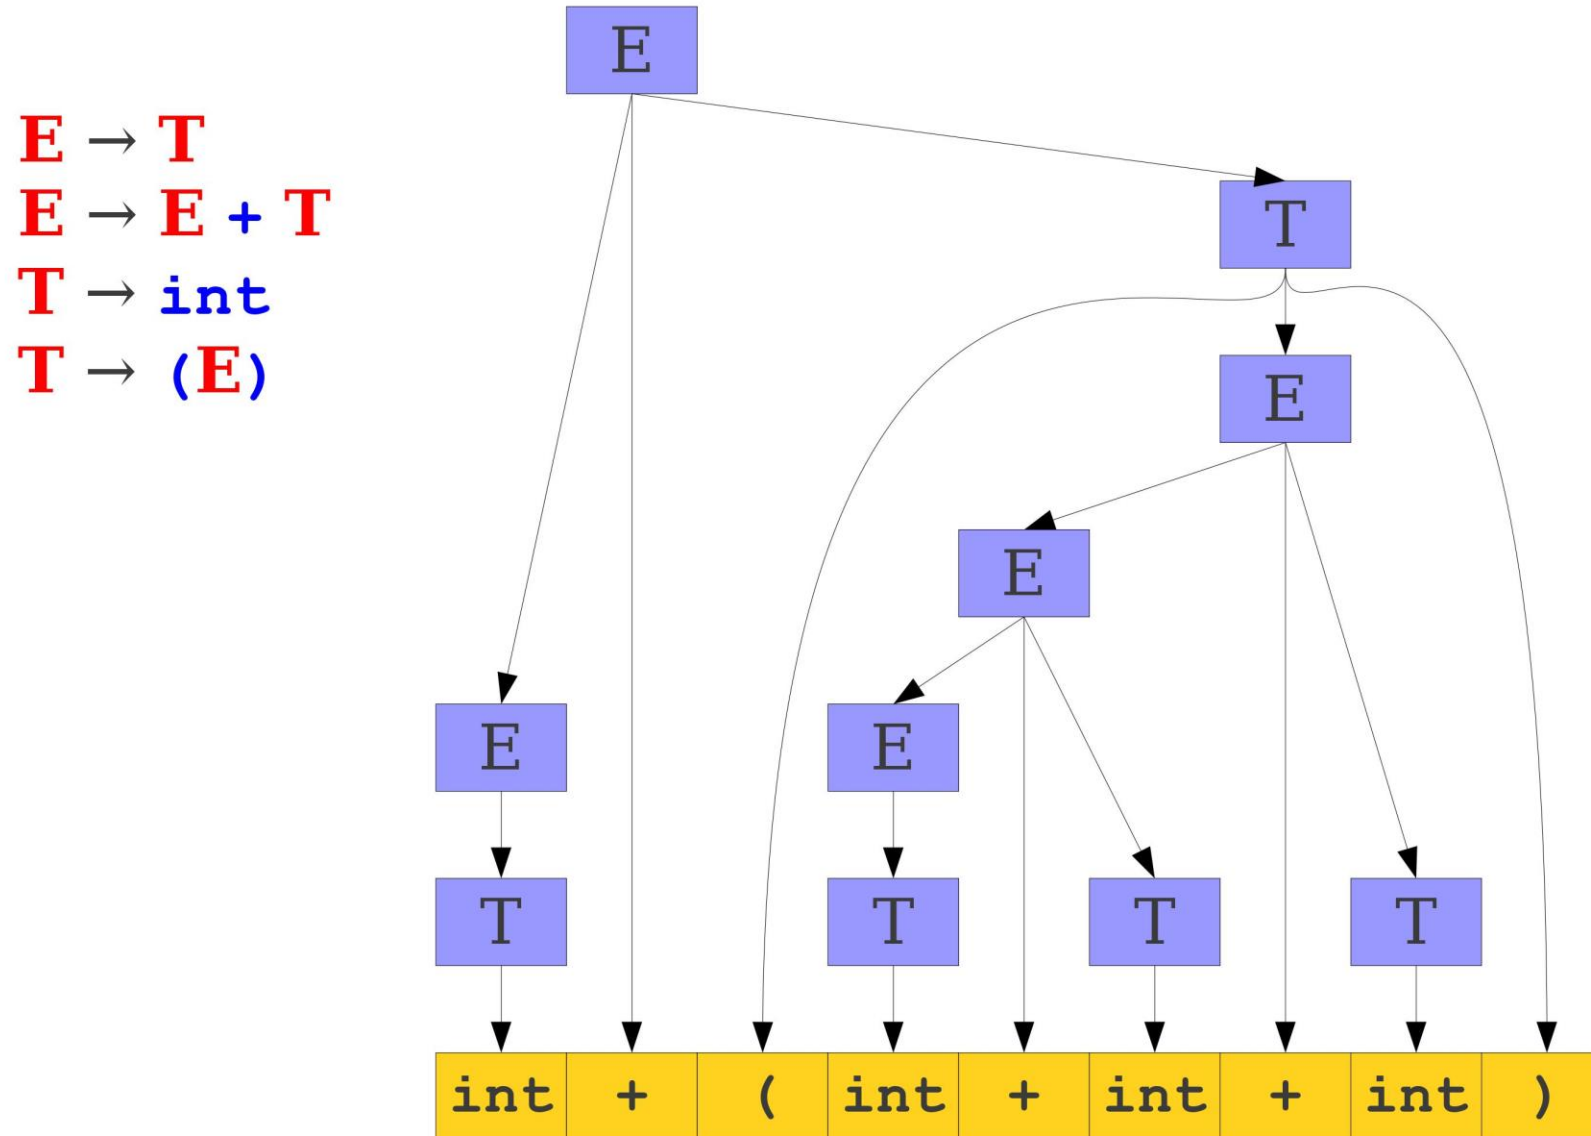$$E \rightarrow E + (E) \mid int$$

- This CFG is this not LL(1).

- If not convinced, consider the string int + (int) + (int).

Introduction example for Bottom-Up Parsing

Procedure, in Layman terms: Bottom-Up Parsers attempts to reduce a string *x* to the start symbol by inverting productions.

- At the beginning, *x* should consist only of terminal symbols.

- Identify a substring β in *x*, such that $A \rightarrow \beta$ is a production rule of our CFG. In other words, it means $x = \alpha\beta\gamma$, with α and γ being strings of some sort (could be empty strings).

- Replace β with A inside of *x*, replacing *x* with $\alpha A\gamma$.

- Keep on doing so, until *x* becomes the start symbol S of the CFG.

Introduction example for Bottom-Up Parsing

$$E \rightarrow T$$
$$E \rightarrow E + T$$
$$T \rightarrow int$$
$$T \rightarrow (E)$$

Introduction example for Bottom-Up Parsing

$E \rightarrow T$
$E \rightarrow E + T$
$T \rightarrow int$
$T \rightarrow (E)$

int + (int + int + int)
$\Rightarrow$ T + (int + int + int)
$\Rightarrow$ E + (int + int + int)
$\Rightarrow$ E + (T + int + int)
$\Rightarrow$ E + (E + int + int)
$\Rightarrow$ E + (E + T + int)
$\Rightarrow$ E + (E + int)
$\Rightarrow$ E + (E + T)
$\Rightarrow$ E + (E)
$\Rightarrow$ E + T
$\Rightarrow$ E

Introduction example for Bottom-Up Parsing

$E \rightarrow T$
$E \rightarrow E + T$
$T \rightarrow \textbf{int}$
$T \rightarrow \textbf{(E)}$

int + (int + int + int)
$\Rightarrow$ T + (int + int + int)
$\Rightarrow$ E + (int + int + int)
$\Rightarrow$ E + (T + int + int)
$\Rightarrow$ E + (E + int + int)
$\Rightarrow$ E + (E + T + int)
$\Rightarrow$ E + (E + int)
$\Rightarrow$ E + (E + T)
$\Rightarrow$ E + (E)
$\Rightarrow$ E + T
$\Rightarrow$ E

Introduction example for Bottom-Up Parsing

Fact #2: A Bottom-Up parser traces a rightmost derivation in reverse.

- This has an interesting consequence…

- Let $\alpha\beta\gamma$ be a step of a bottom-up parse.

- Assume the next reduction is by $A \rightarrow \beta$

- **Then γ is necessarily a string of terminals!**

- Why? Because $\alpha A\gamma \rightarrow \alpha\beta\gamma$ is a step in a rightmost derivation!
  (If γ does not change, it means it does not contain any non-terminals)

Introduction example for Bottom-Up Parsing

int + (int + int + int)
⇒ **T** + (int + int + int)
⇒ **E** + (int + int + int)
⇒ **E** + (**T** + int + int)
⇒ **E** + (**E** + int + int)
⇒ **E** + (**E** + **T** + int)
⇒ **E** + (**E** + int)
⇒ **E** + (**E** + **T**)
⇒ **E** + (**E**)
⇒ **E** + **T**
⇒ **E**
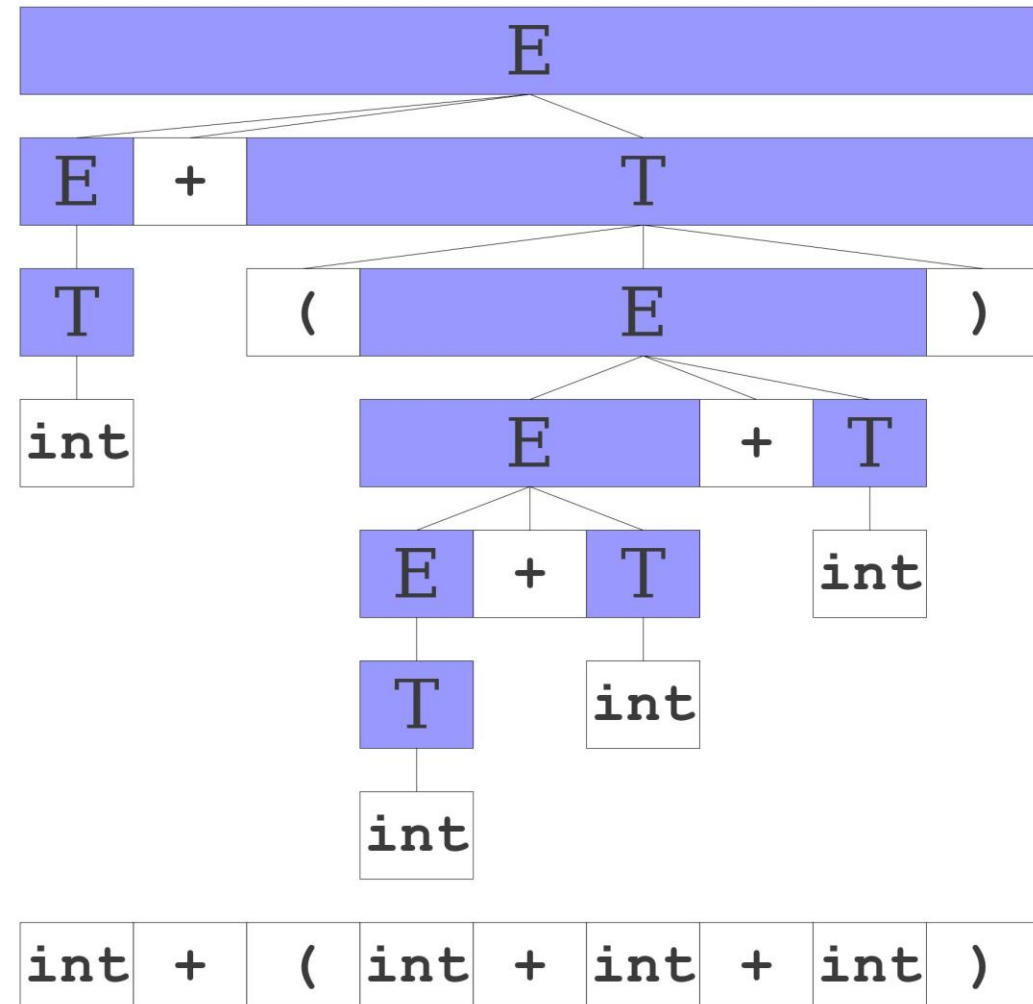
Introduction example for Bottom-Up Parsing

Follow-up idea from Fact #2: Split the string into two substrings.

- The right substring should consist only of terminal symbols, and has yet to be examined by the parser.

- The left substring could have terminals and non-terminals.

- Mark the dividing point using the | symbol.
  (Note that this symbol | is only for visualization purposes, it is not part of the string to be analysed!)

- At the beginning, the string x is therefore written as $x = |x_1 x_2 \ldots x_n$\$, with $x_1, x_2, \ldots x_n$ being terminal symbols.

Shift-Reduce Parsing

Fact #3: A Bottom-up Parser will attempt to revert the string by using only two possible actions.

- **Shifting:** Moves the separator one step to the right (one full symbol).

For instance, E|+ (int) becomes E +|(int) after shifting.

- **Reducing:** Apply an inverse production rule at the right of the left end string.

For instance, the string E + (E + (E)|) can be reduced into E + (E |) by using the production E → E + (E).

Practice 1: Shift-Reduce tryout

Question: Consider the CFG on the right, what is the correct sequence of Reduce and Shift operations to use on the string *x* below?
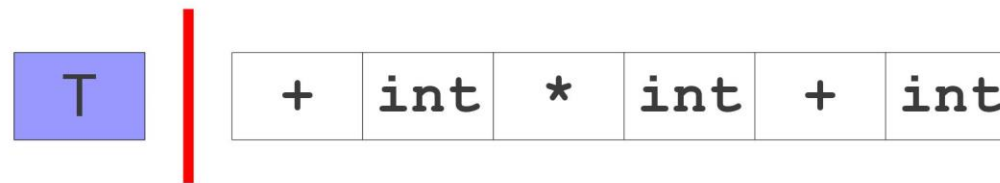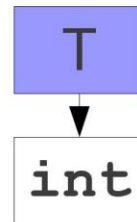
int + int * int + int

$$E \rightarrow F$$
$$E \rightarrow E + F$$
$$F \rightarrow F * T$$
$$F \rightarrow T$$
$$T \rightarrow \text{int}$$
$$T \rightarrow (E)$$

Practice 1: Shift-Reduce tryout

E → F
E → E + F
F → F * T
F → T
T → int
T → (E)

| int | + | int | * | int | + | int |

Practice 1: Shift-Reduce tryout

E → F
E → E + F
F → F * T
F → T
T → int
T → (E)

| int | | + | int | * | int | + | int |

Practice 1: Shift-Reduce tryout

$E \to F$
$E \to E + F$
$F \to F * T$
$F \to T$
$T \to int$
$T \to (E)$

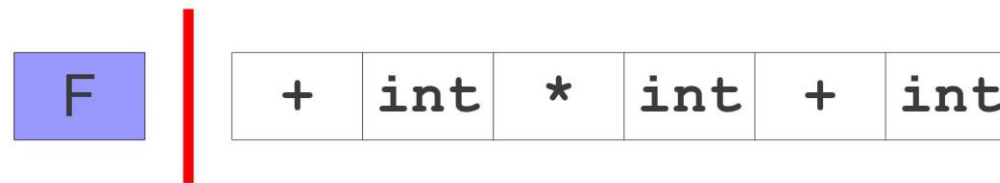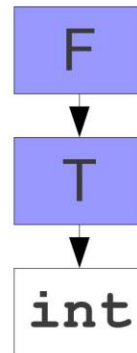# Practice 1: Shift-Reduce tryout

$E \rightarrow F$
$E \rightarrow E + F$
$F \rightarrow F * T$
$F \rightarrow T$
$T \rightarrow int$
$T \rightarrow (E)$

Practice 1: Shift-Reduce tryout

$E \rightarrow F$
$E \rightarrow E + F$
$F \rightarrow F * T$
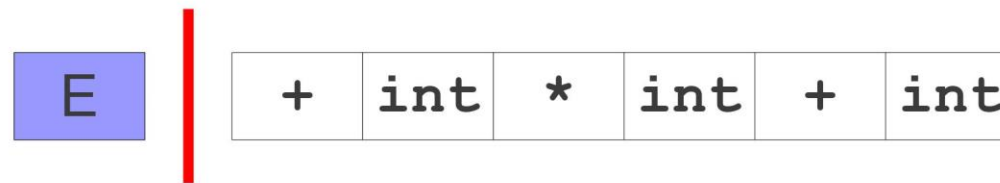$F \rightarrow T$
$T \rightarrow int$
$T \rightarrow (E)$

Practice 1: Shift-Reduce tryout

$E \rightarrow F$
$E \rightarrow E + F$
$F \rightarrow F * T$
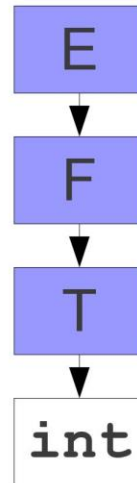$F \rightarrow T$
$T \rightarrow int$
$T \rightarrow (E)$

Practice 1: Shift-Reduce tryout

$E \rightarrow F$
$E \rightarrow E + F$
$F \rightarrow F * T$
$F \rightarrow T$
$T \rightarrow int$
$T \rightarrow (E)$

Practice 1: Shift-Reduce tryout

$E \rightarrow F$
$E \rightarrow E + F$
$F \rightarrow F * T$
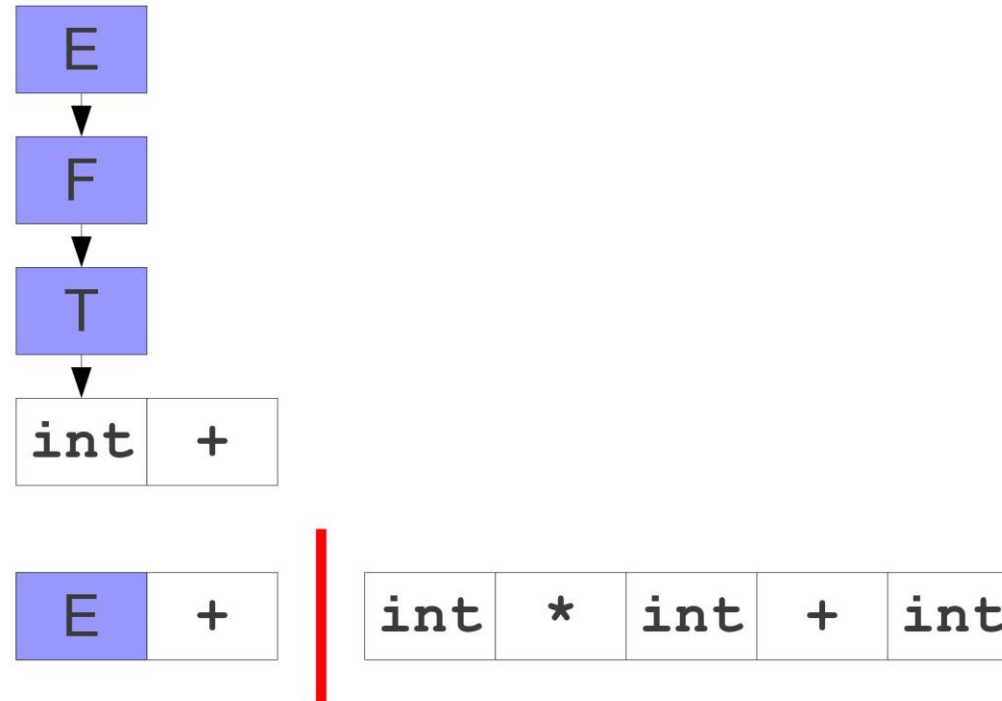$F \rightarrow T$
$T \rightarrow int$
$T \rightarrow (E)$

Practice 1: Shift-Reduce tryout

$E \rightarrow F$
$E \rightarrow E + F$
$F \rightarrow F * T$
$F \rightarrow T$
$T \rightarrow int$
$T \rightarrow (E)$

Practice 1: Shift-Reduce tryout

$E \rightarrow F$
$E \rightarrow E + F$
$F \rightarrow F * T$
$F \rightarrow T$
$T \rightarrow int$
$T \rightarrow (E)$

Practice 1: Shift-Reduce tryout

$E \rightarrow F$
$E \rightarrow E + F$
$F \rightarrow F * T$
$F \rightarrow T$
$T \rightarrow int$
$T \rightarrow (E)$

Practice 1: Shift-Reduce tryout

E → F
E → E + F
F → F * T
F → T
T → int
T → (E)

Practice 1: Shift-Reduce tryout

E → F
E → E + F
F → F * T
F → T
T → int
T → (E)

# Practice 1: Shift-Reduce tryout

$E \rightarrow F$
$E \rightarrow E + F$
$F \rightarrow F * T$
$F \rightarrow T$
$T \rightarrow int$
$T \rightarrow (E)$

Practice 1: Shift-Reduce tryout

Practice 1: Shift-Reduce tryout

$E \rightarrow F$
$E \rightarrow E + F$
$F \rightarrow F * T$
$F \rightarrow T$
$T \rightarrow int$
$T \rightarrow (E)$

Practice 1: Shift-Reduce tryout

$E \to F$
$E \to E + F$
$F \to F * T$
$F \to T$
$T \to int$
$T \to (E)$

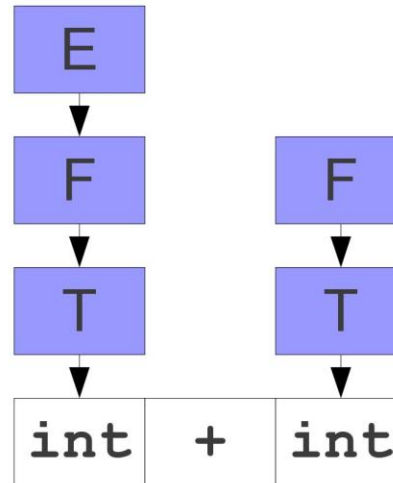# Practice 1: Shift-Reduce tryout

$E \rightarrow F$
$E \rightarrow E + F$
$F \rightarrow F * T$
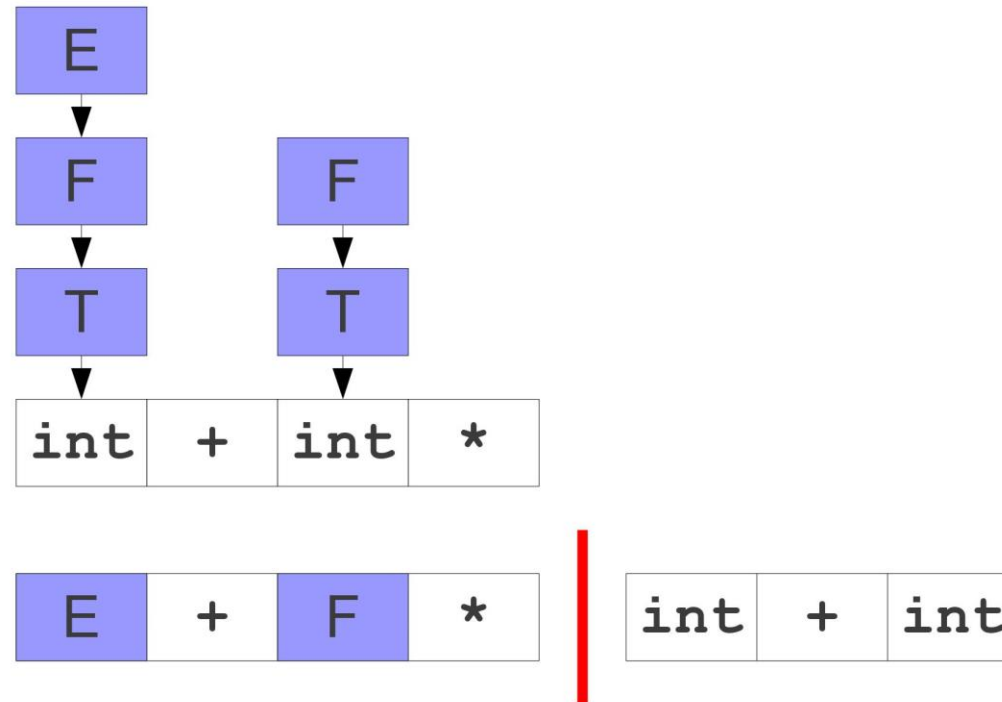$F \rightarrow T$
$T \rightarrow int$
$T \rightarrow (E)$

Practice 1: Shift-Reduce tryout

E → F
E → E + F
F → F * T
F → T
T → int
T → (E)

Practical implementation of Shift-Reduce

Practical implementation for the Shift-Reduce parser can be done as follows.

- **Left substring** can be implemented by a **stack**.

- Top of the stack is denoted by the separator symbol **|**.

- **Shifting pushes a terminal symbol on the stack.**

- **Reducing using a production rule A → B pops all the symbols in B off the top of the stack and then pushes a non-terminal symbol A on the stack to replace all the symbols in B.**

- (Also, two more actions technically being: **accept** the string if parsing has completed, or **reject** the string if derivation is not possible due to invalid syntax...)

Million dollar question

Million dollar question: How to decide when to Shift, when to Reduce, and with which production rule from the CFG?

A few observations we can make so far

1. **You can always shift at any time,** except when the separator | has reached the end of the string.

Million dollar question

Million dollar question: How to decide when to Shift, when to Reduce, and with which production rule from the CFG?

A few observations we can make so far

1. **You can always shift at any time,** except when the separator | has reached the end of the string.

2. **When you cannot reduce the string using any production rule of the CFG, shift, by default.**

Million dollar question

Million dollar question: How to decide when to Shift, when to Reduce, and with which production rule from the CFG?

A few observations we can make so far

1. **You can always shift at any time,** except when the separator | has reached the end of the string.

2. **When you cannot reduce the string using any production rule of the CFG, shift, by default.**

3. **The difficulty comes from deciding between shifting and reducing in scenarios where you could do both. To decide, we will reuse the idea of LL(1), about looking at the first terminal symbol to appear after the split symbol |.**

How to decide between shift/reduce?

First Idea: Build an FSM (again!), to decide when to shift or reduce.

- FSM will use, as input string, the current stack.
- The **actions** consist of the possible **terminals and non-terminals symbols** that the CFG possesses.

How to decide between shift/reduce?

First Idea: Build an FSM (again!), to decide when to shift or reduce.

- The **states** will have **shift-reduce(rule)-accept-reject actions** assigned to them.

- Whatever final state we end up on after running the FSM will tell us which parsing action between shift and reduce.

- In the case of a reduce action, we will also mention the production rule to use to transform elements in our stack.

- Also, keep in mind that this is an LR(1) parser. This means one lookahead symbol should be used. In our case, this means that **the parsing action to use for a given final state might depend on what is currently the first terminal symbol in the right substring, after |.**

An example FSM for our CFG

Consider our simplified CFG from earlier, below.

$$1: E \rightarrow E + (E)$$

$$2: E \rightarrow int$$

Our FSM has the following **actions**: E, (,), +, int, $.

The parser has the following **parsing actions** to be assigned to **states**: shift, reduce using rule 1, reduce using rule 2, accept, reject.

For simplicity, we will not draw all **states transitions**.

If a state transition is used, but does not appear in the FSM diagram (on the next slide), then we should stop the FSM and reject the string due to invalid syntax!

An example FSM for our CFG

Consider our simplified CFG from earlier, below.

$$1: E \rightarrow E + (E)$$

$$2: E \rightarrow int$$

The LR(1) FSM hiding behind the CFG above is then…
(Given on the next slide)
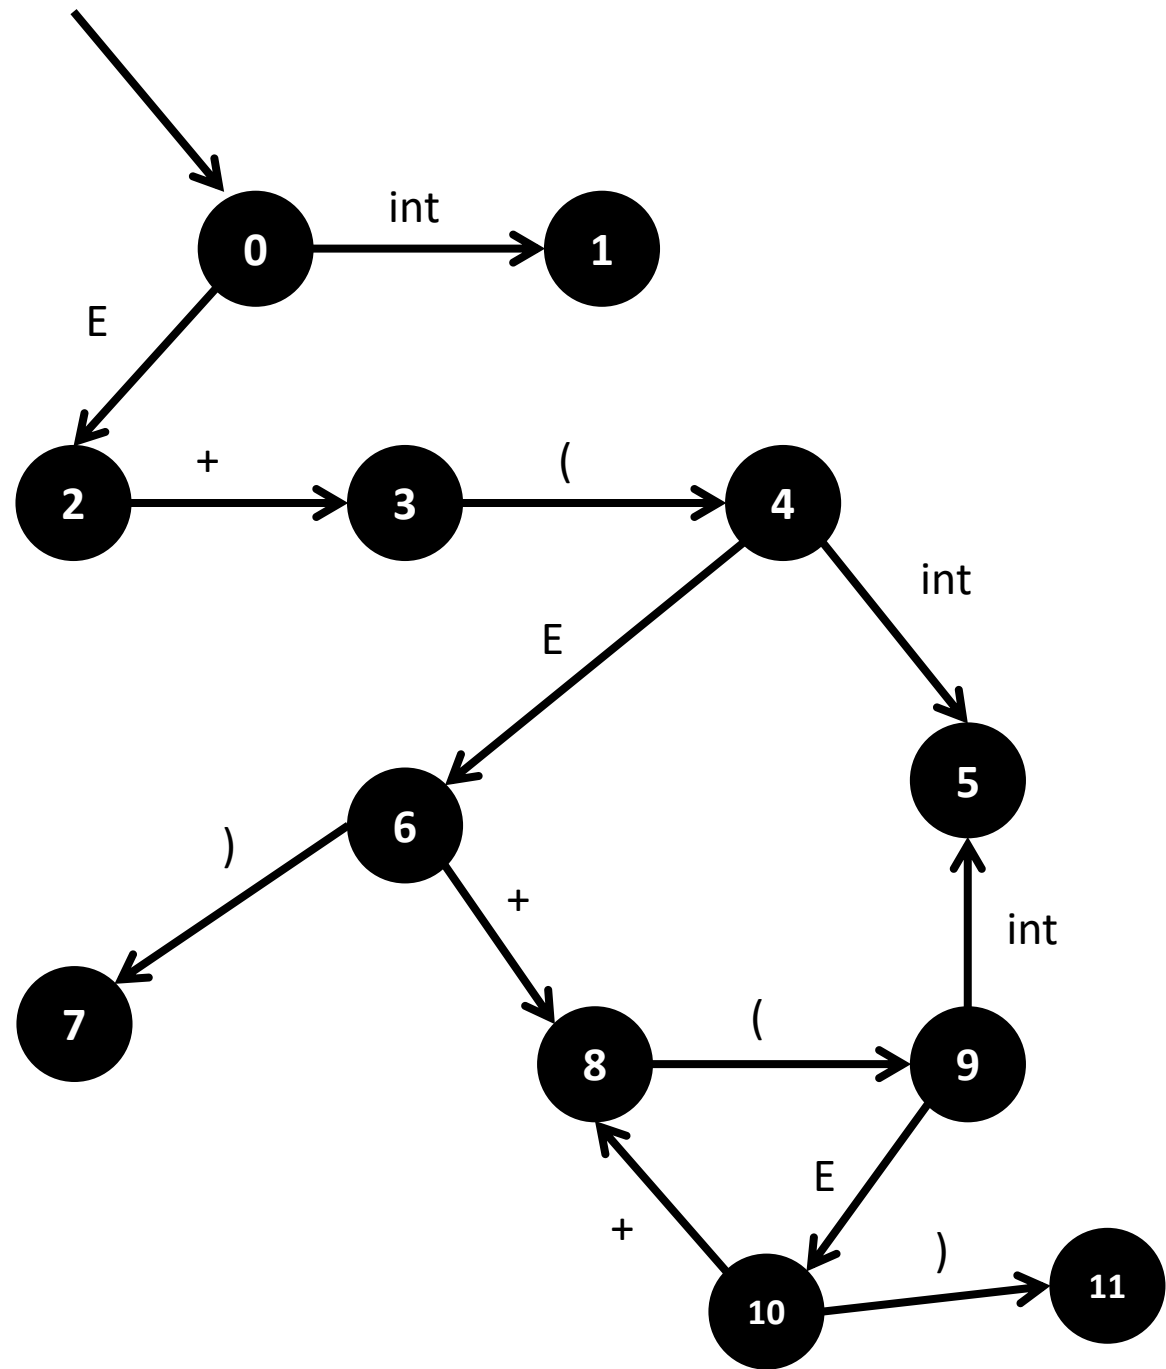
State	Parsing action
0	Shift
1	Reduce using 2 on $ or +
2	Accept on $
3	Shift
4	Shift
5	Reduce using 2 on) or +
6	Shift
7	Reduce using 1 on $ or +
8	Shift
9	Shift
10	Shift
11	Reduce using 1 on) or +

Note: shift by default, if the next symbol after | is not listed. For instance, on state 1, we use reduce if the element after | is $ or +; and shift otherwise.
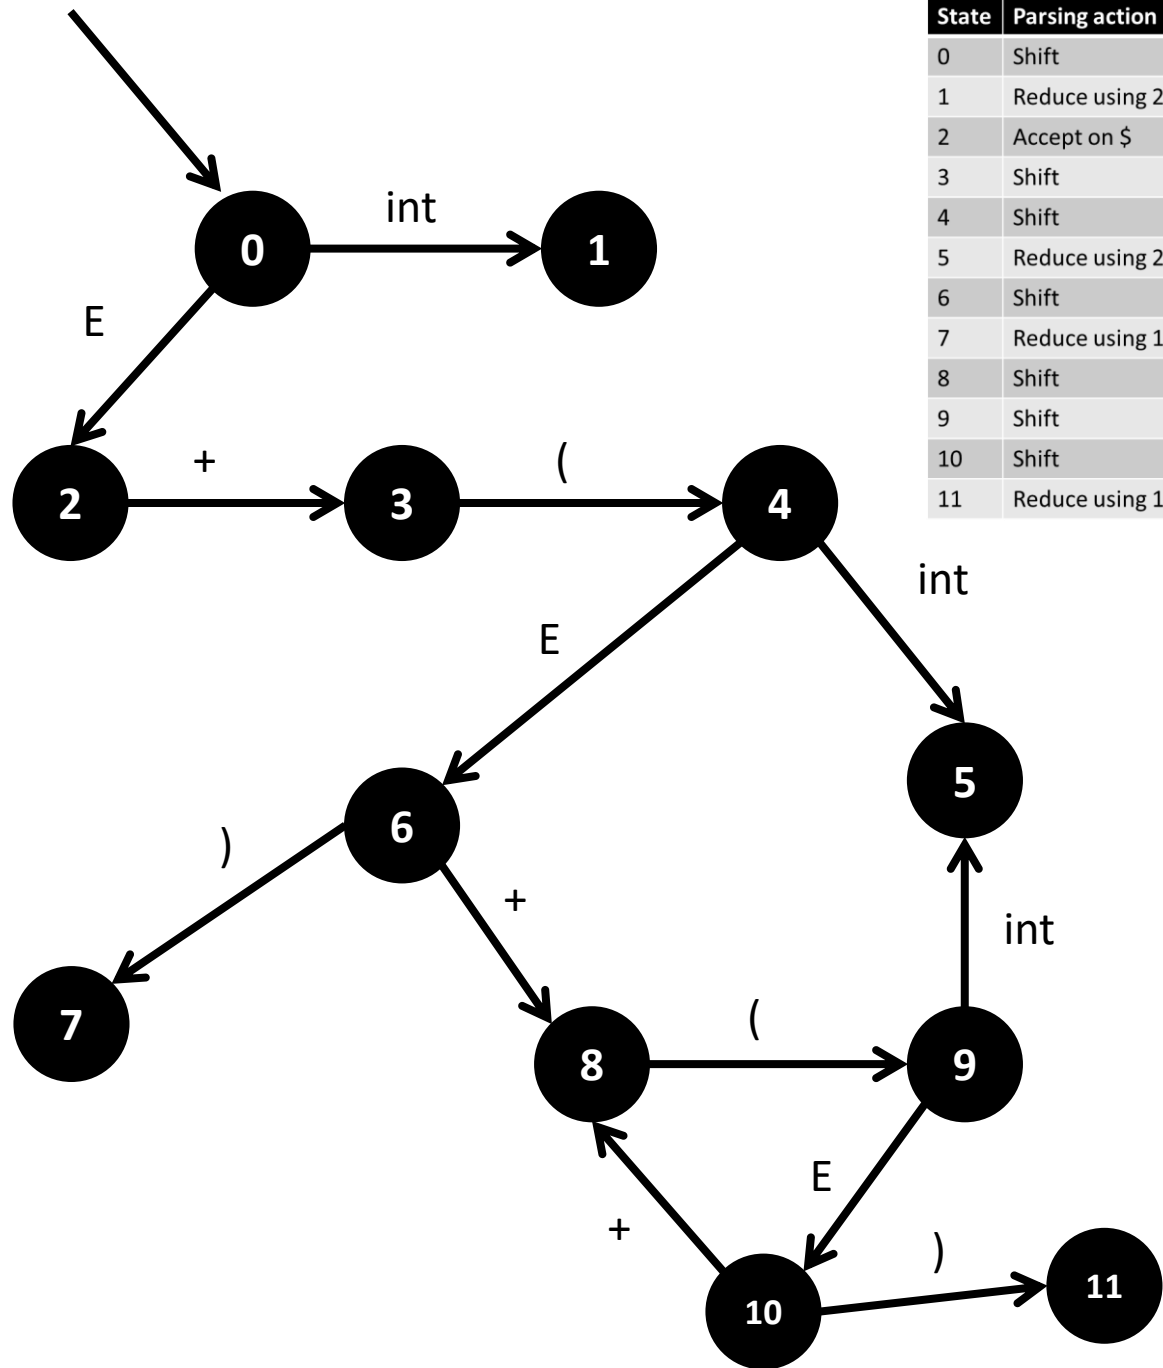
State	Parsing action
0	Shift
1	Reduce using 2 on $ or +
2	Accept on $
3	Shift
4	Shift
5	Reduce using 2 on) or +
6	Shift
7	Reduce using 1 on $ or +
8	Shift
9	Shift
10	Shift
11	Reduce using 1 on) or +

Let us try it on the string int + (int) + (int)!

| int + (int) + (int)$	[0, shift]
int | + (int) + (int)$	[01, reduce using 2]
E | + (int) + (int)$	[02, shift]
E +| (int) + (int)$	[023, shift]
E + (|int) + (int)$	[0234, shift]
E + (int|) + (int)$	[02345, reduce using 2]
E + (E|) + (int)$	[02346, shift]
E + (E)| + (int)$	[023467, reduce using 1]
E| + (int)$	[02, shift]
E+| (int)$	[023, shift]
E+(|int)$	[0234, shift]
E+(int|)$	[02345, reduce using 1]
E+(E|)$	[02346, shift]
E+(E)|$	[023467, reduce using 1]
E|$	[02, accept]

That works, first try, and did not have to branch and backtrack either!

That is very powerful!

- This parser, whose name would be LR(1) according to our naming convention from earlier,...

- ...Definitely works!

Several follow-up questions, obviously:

1. How did we come up with the FSM and the FSM state to parsing actions mapping for our LR(1) parser?

2. Does the CFG need to have special properties, again, for this to work?

Addressing the questions

Does the CFG need to have special properties, again, for this to work?

- Yes, the CFG needs to have certain properties, which will appear when we start building the FSM following from that CFG.

- We say that the **CFG is LR(1)**, if such an LR(1) parser can be used on it.

- **Good news #1:** LR(1) is a stronger type than LL(1). This means our LR(1) parser can cover more CFGs than LL(1)!

Addressing the questions

LR(1) is a stronger type than LL(1). This means our LR(1) parser can cover more CFGs than LL(1)!

Also, did not mention, but LL(k) and LR(k) grammars are non-ambiguous!

Ambiguous grammars are problematic (refer to W10S2!).

Addressing the questions

Does the CFG need to have special properties, again, for this to work?

- Yes, the CFG needs to have certain properties, which will appear when we start building the FSM following from that CFG.

- We say that the **CFG is LR(1)**, if such an LR(1) parser can be used on it.

- **Good news #1:** LR(1) is a stronger type than LL(1). This means our LR(1) parser can cover more CFGs than LL(1)!

- **Good news #2:** LR(1) is a very large type for CFGs, and most programming languages CFGs are LR(1)!
  (For the other programming languages that are not LR(1), this means that more advanced parsers are required!)

Addressing the questions

But how did we come up with the FSM and the FSM states to parsing actions mapping for our LR(1) parser?

- This will be the not-so-easy part of figuring out the FSM hiding behind a given CFG!

- Typically done by **parsing generators**.

- A parsing generator will look at a CFG and build a non-deterministic FSM first, eventually reducing it to a deterministic FSM later.

Production tracking in production rules

Definition (The point symbol in LR(1)):

The **point symbol** **(.)** in the context of LR(1) parsing is used to represent the **current position in a production rule while parsing a given input**.

It helps us **keep track of how much of the production rule has been parsed and what still needs to be parsed.**

In simple terms, the dot (.) separates the part of the production rule that has been recognized (to the left of the dot) from the part that still needs to be parsed and seen via shifting (to the right of the dot).

Production tracking in production rules

For example, consider the production rule A → BCD.

- An LR(1) item with the dot at the beginning of the rule, represented as A → .BCD

- This indicates that none of the right-hand side of the production rule has been recognized yet.

As the parser progresses and recognizes the input, the dot moves to the right:

- A → B.CD   (after recognizing B)

- A → BC.D   (after recognizing C)

- A → BCD.   (after recognizing D)

Production tracking in production rules

- When the dot is at the beginning or in the middle of a production rule (i.e. A → .BCD, A → B.CD, or A → BC.D), **we have not yet recognized enough symbols to be sure we can reduce BCD into A**.

- When the dot reaches the end of the production rule (i.e. A -> BCD.), it means the entire rule has been recognized, and the **parser could technically chose to perform a reduce action based on that production rule**.

Production tracking in production rules

An important theorem:

- Any production rule necessarily has a finite number of elements on the right hand side. Technically, this means that **there is a limited number of possibilities for the dot to be at** for any given production rule.

For instance, the production rule A → BCD has four possibilities in terms of dot positions, being:

- *A → .BCD*

- *A → B.CD*

- *A → BC.D*

- *A → BCD.*

Production tracking in production rules

An important theorem:

- Any production rule necessarily has a finite number of elements on the right hand side. Technically, this means that **there is a limited number of possibilities for the dot to be at** for any given production rule.

- But **our CFGs also have a finite number of production rules**.

For instance, the CFG below has four possible production rules.

$$E \rightarrow T,$$
$$E \rightarrow E + T,$$
$$T \rightarrow (E),$$
$$T \rightarrow int$$

Production tracking in production rules

An important theorem:

- Any production rule necessarily has a finite number of elements on the right hand side. Technically, this means that **there is a limited number of possibilities for the dot to be at** for any given production rule.

- But **our CFGs also have a finite number of production rules**.

- Combining these two pieces of information together means that **there is a finite number of possible dot representation for all the production rules of our CFG**.

- And these could be used as **states for an FSM of some sort!**

A finite number of dot representations

For instance, the CFG below has four possible production rules.

$$E \rightarrow T,$$

$$E \rightarrow E + T,$$

$$T \rightarrow (E),$$

$$T \rightarrow int$$

According to our theorem, it means that there should be a finite number of dot representations for the production rules of that CFG.

That is true and shown below.

(Note: we have also added the dot representations for the start symbol of our CFG, which is E).

- .E, E.,
- .T, T.,
- .E+T, E.+T, E+.T, E+T.,
- .(E), (.E), (E.), (E).,
- .int, int.

From LR(0) to LR(1) items

Definition (LR(0) items for a given CFG):

For a given CFG, **all the possible right hand side formulas of all the production rules and their possible variations in terms of dot positions**, produce what is commonly referred to as the **LR(0) items** of the CFG. In the case of our previous CFG, the LR(0) items are simply:

- .E, E.,
- .T, T.,
- .E+T, E.+T, E+.T, E+T.,
- .(E), (.E), (E.), (E).,
- .int, int.

From LR(0) to LR(1) items

Definition (LR(1) items for a given CFG):

For a given CFG, the LR(1) items are the set of all possible tuples (A → α.Bβ, a), where

- A is a non-terminal symbol,

- α and β are strings of terminal and nonterminal symbols,

- B is a nonterminal symbol,

- and a is a lookahead symbol (all possible terminal symbols in the CFG or the end of string $ symbol).

From LR(0) to LR(1) items

Some examples of LR(1) items for our CFG are then:

- **(E → .T, +):** This LR(1) item indicates that we are starting to parse the non-terminal E, and we expect to see a T followed by a + symbol.

- **(E → .T,)):** This LR(1) item indicates that we are parsing the non-terminal E, and we expect to see a T followed by a closing parenthesis.

- **(E → E +. T, +):** This LR(1) item indicates that we are parsing the non-terminal E, and we have already seen some prefix of the form E + ... and we expect to see a T followed by a + symbol.

Important: there is a finite number of LR(1) items in any given CFG.

Building an FSM: states

- Let us use these LR(1) items as **states for an FSM of some sort!**

- How about the **input string** to be used for that FSM?
  As before, let us use the current left string in our parser stack.

- How about our FSM **actions**? As before, could consist of any terminal or non-terminal symbols in our CFG.

- How to establish a **transition logic** for our FSM then?

Transition logic rules: Rule #1

Rule #1, Shift: If there is an LR(1) item of the form (A $\rightarrow \alpha.X\beta$, s) used as the current state and X **is a terminal symbol**, add a transition to the state containing the item (A $\rightarrow \alpha X.\beta$, s) using the symbol X.

Transition logic rules: Rule #1

Etc.

Transition logic rules: Rule #1

Rule #1, Shift: If there is an LR(1) item of the form (A → α.Xβ, s) used as the current state and X **is a terminal symbol**, add a transition to the state containing the item (A → αX.β, s) using the symbol X.

Rule #2, GoTo: If there is an LR(1) item of the form (A → α.Xβ, s) used as the current state and X **is a non-terminal**, add a transition to the state containing the item (A → αX.β, s) on the symbol X.
(Roughly similar)

Transition logic rules: Rule #2

Etc.

Transition logic rules: Rule #3

Rule #3, Closure: If there is an LR(1) item of the form (A → α.Xβ, s) used as the current state and X **is a non-terminal**, for each production X → γ and each terminal a in FIRST*(βs), add the item (X → .γ, a) to the current state if it is not already there.

This means that some states will combine together.

Transition logic rules: Rule #4

Rule #3, Closure: If there is an LR(1) item of the form (A → α.Xβ, s) used as the current state and X **is a non-terminal**, for each production X → γ and each terminal a in FIRST*(βs), add the item (X → .γ, a) to the current state if it is not already there.

This means that some states will combine together.

Rule #4, Start State: The start state of the LR(1) parsing automaton contains the item (S' → . E, $), where S' is the augmented start symbol and $ is the end-of-input marker.

Tedious, not fun, but technically possible
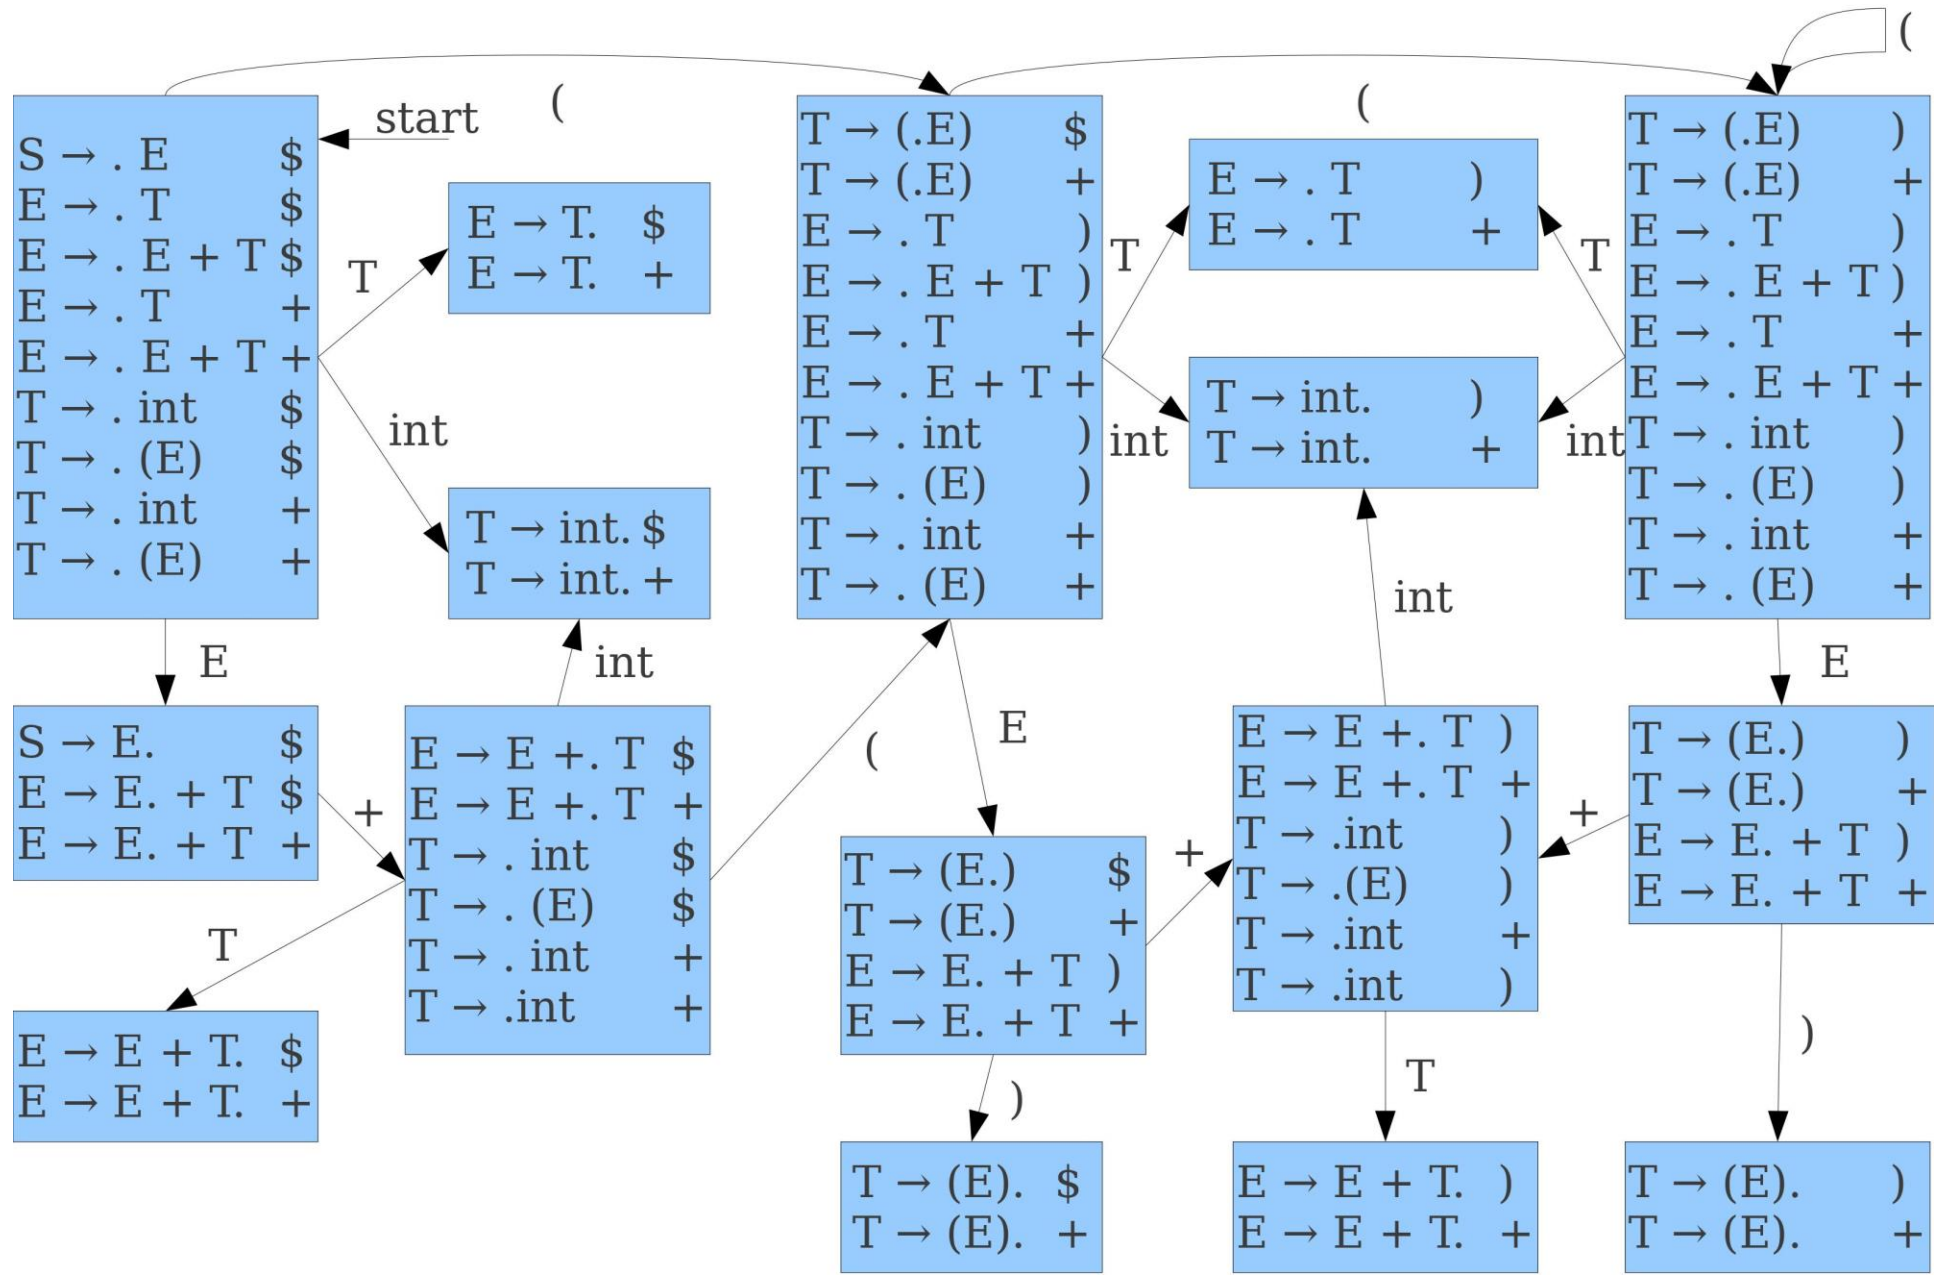
Getting there!

Ok, so now, it means that we have managed to produce a (massive) **deterministic FSM.**

- This FSM is called the LR(1) finite state automata for the CFG and will be used by the LR(1) parser, in an almost identical way as shown earlier.

- The LR(1) will now run as follows, using one of four possible actions.

Lifting the final shift-reduce conflict

The LR(1) will now run as follows, using one of four possible actions.

1. **Shift**: If the parser is in a state containing an LR(1) item of the form (A → α.Xβ, s) and the next input symbol is X (a terminal symbol), the parser will shift the input symbol onto the stack and move to the state connected by the X transition.

Lifting the final shift-reduce conflict

The LR(1) will now run as follows, using one of four possible actions.

2. **Reduce**: If the parser is in a state containing an LR(1) item of the form (A → α. , s) and the next input symbol matches the lookahead symbol s, the parser will perform a reduction using the production rule A → α. It will pop|α|symbols (|α|being the length of α) from the stack, go to the state at the top of the stack, follow the GOTO transition on A to the next state, and push that state onto the stack.

Lifting the final shift-reduce conflict

The LR(1) will now run as follows, using one of four possible actions.

3. **Accept**: If the parser is in a state containing an LR(1) item of the form (S' → E. , $) and the next input symbol is $ (end-of-input marker), the parser will accept the input string as valid.

Lifting the final shift-reduce conflict

The LR(1) will now run as follows, using one of four possible actions.

4. **Error**: If the parser encounters an input symbol that does not correspond to any valid shift or reduce action based on the current state, it will report a parsing/syntax error.

To summarize

- While parsing an input string, the LR(1) parser will iteratively perform these actions based on the parsing automaton's states and transitions until it either accepts the input, encounters an error, or exhausts the input string.

- In summary, the LR(1) parser will decide when to shift or reduce based on the LR(1) items in the current state and the next input symbol.

- It will shift when the next input symbol matches a terminal symbol in an LR(1) item with a dot before it.

- And it will reduce when the next input symbol matches the lookahead symbol in an LR(1) item with a dot at the end.

But wait…

According to our previous definitions, LR(1) will always attempt to reduce whenever possible.

But, technically, it feels like could decide to reduce, or I could decide to shift some more symbols and reduce later…!

I feel that this could be a problem in certain CFGs!

Question: Are there scenarios where we should postpone an opportunity to reduce, shift some more, and then reduce using a different production rule?

A quick word about conflicts

Definition (Shift-Reduce conflicts):

A **shift-reduce conflict** occurs in a parsing situation **when the parser cannot decide whether to shift the next input symbol or reduce by applying a production rule**.

This **ambiguity** can lead to different parse trees for the same (syntactically valid) input string.

Conflicts may arise in grammars that are not LR(1), meaning they cannot be unambiguously parsed using an LR(1) parser.

In those cases, more advanced parsers, like LR(k), must be used (out-of-scope though!)

A quick word about conflicts

Definition (Reduce-Reduce conflicts):

A **reduce-reduce conflict** occurs **when a parser, in a single state, has two or more reduction actions available with different production rules for the same lookahead symbol**.

It means that the parser cannot unambiguously decide which production rule to apply for reducing the current input.

Like shift-reduce conflicts, reduce-reduce conflicts may arise in grammars that are not LR(1).

In those cases, more advanced parsers, like LR(k), must be used (out-of-scope though!)

But wait…

According to our previous definitions, LR(1) will always attempt to reduce whenever possible.

But, technically, it feels like could decide to reduce, or I could decide to shift some more symbols and reduce later…!

I feel that this could be a problem in certain CFGs!

Question: Are there scenarios where we should postpone an opportunity to reduce, shift some more, and then reduce using a different production rule?

Answer: If the CFG is LR(1), no need to worry about that, reduce when you can and no conflicts will arise.

An important question

Erm, will sound like a stupid question, but is our CFG actually LR(1)?

$$E \rightarrow T,$$

$$E \rightarrow E + T,$$

$$T \rightarrow (E),$$

$$T \rightarrow int$$

Yes, it is!

An important question

Ok, how about the YACC CFG that you have shown us for the entirety of the C programming language in a previous lecture?

(http://www.lysator.liu.se/c/ANSI-C-grammar-y.html)

...

```
%token IDENTIFIER CONSTANT STRING_LITERAL SIZEOF
%token PTR_OP INC_OP DEC_OP LEFT_OP RIGHT_OP LE_OP GE_OP EQ_OP NE_OP
%token AND_OP OR_OP MUL_ASSIGN DIV_ASSIGN MOD_ASSIGN ADD_ASSIGN
%token SUB_ASSIGN LEFT_ASSIGN RIGHT_ASSIGN AND_ASSIGN
%token XOR_ASSIGN OR_ASSIGN TYPE_NAME

%token TYPEDEF EXTERN STATIC AUTO REGISTER
%token CHAR SHORT INT LONG SIGNED UNSIGNED FLOAT DOUBLE CONST VOLATILE VOID
%token STRUCT UNION ENUM ELLIPSIS

%token CASE DEFAULT IF ELSE SWITCH WHILE DO FOR GOTO CONTINUE BREAK RETURN

%start translation_unit
%%

primary_expression
        : IDENTIFIER
        | CONSTANT
        | STRING_LITERAL
        | '(' expression ')'
        ;

postfix_expression
        : primary_expression
        | postfix_expression '[' expression ']'
        | postfix_expression '(' ')'
        | postfix_expression '(' argument_expression_list ')'
        | postfix_expression '.' IDENTIFIER
        | postfix_expression PTR_OP IDENTIFIER
        | postfix_expression INC_OP
        | postfix_expression DEC_OP
        ;

argument_expression_list
        : assignment_expression
        | argument_expression_list ',' assignment_expression
        ;

unary_expression
        : postfix_expression
        | INC_OP unary_expression
        | DEC_OP unary_expression
        | unary_operator cast_expression
        | SIZEOF unary_expression
        | SIZEOF '(' type_name ')'
        ;

unary_operator
        : '&'
        | '*'
        | '+'
        | '-'
        | '~'
        | '!'
        ;

cast_expression
        : unary_expression
        | '(' type_name ')' cast_expression
        :
```

An important question

Ok, how about the YACC CFG that you have shown us for the entirety of the C programming language in a previous lecture?

(http://www.lysator.liu.se/c/ANSI-C-grammar-y.html)

Yes, it is!

```
%token IDENTIFIER CONSTANT STRING_LITERAL SIZEOF
%token PTR_OP INC_OP DEC_OP LEFT_OP RIGHT_OP LE_OP GE_OP EQ_OP NE_OP
%token AND_OP OR_OP MUL_ASSIGN DIV_ASSIGN MOD_ASSIGN ADD_ASSIGN
%token SUB_ASSIGN LEFT_ASSIGN RIGHT_ASSIGN AND_ASSIGN
%token XOR_ASSIGN OR_ASSIGN TYPE_NAME

%token TYPEDEF EXTERN STATIC AUTO REGISTER
%token CHAR SHORT INT LONG SIGNED UNSIGNED FLOAT DOUBLE CONST VOLATILE VOID
%token STRUCT UNION ENUM ELLIPSIS

%token CASE DEFAULT IF ELSE SWITCH WHILE DO FOR GOTO CONTINUE BREAK RETURN

%start translation_unit
%%

primary_expression
        : IDENTIFIER
        | CONSTANT
        | STRING_LITERAL
        | '(' expression ')'
        ;

postfix_expression
        : primary_expression
        | postfix_expression '[' expression ']'
        | postfix_expression '(' ')'
        | postfix_expression '(' argument_expression_list ')'
        | postfix_expression '.' IDENTIFIER
        | postfix_expression PTR_OP IDENTIFIER
        | postfix_expression INC_OP
        | postfix_expression DEC_OP
        ;

argument_expression_list
        : assignment_expression
        | argument_expression_list ',' assignment_expression
        ;

unary_expression
        : postfix_expression
        | INC_OP unary_expression
        | DEC_OP unary_expression
        | unary_operator cast_expression
        | SIZEOF unary_expression
        | SIZEOF '(' type_name ')'
        ;

unary_operator
        : '&'
        | '*'
        | '+'
        | '-'
        | '~'
        | '!'
        ;

cast_expression
        : unary_expression
        | '(' type_name ')' cast_expression
```

An important question

Ok, how about the YACC CFG that you have shown us for the entirety of the C programming language in a previous lecture?

(http://www.lysator.liu.se/c/ANSI-C-grammar-y.html)

Yes, it is!

This means an LR(1) parser will work nicely for a C compiler!

```
%token IDENTIFIER CONSTANT STRING_LITERAL SIZEOF
%token PTR_OP INC_OP DEC_OP LEFT_OP RIGHT_OP LE_OP GE_OP EQ_OP NE_OP
%token AND_OP OR_OP MUL_ASSIGN DIV_ASSIGN MOD_ASSIGN ADD_ASSIGN
%token SUB_ASSIGN LEFT_ASSIGN RIGHT_ASSIGN AND_ASSIGN
%token XOR_ASSIGN OR_ASSIGN TYPE_NAME

%token TYPEDEF EXTERN STATIC AUTO REGISTER
%token CHAR SHORT INT LONG SIGNED UNSIGNED FLOAT DOUBLE CONST VOLATILE VOID
%token STRUCT UNION ENUM ELLIPSIS

%token CASE DEFAULT IF ELSE SWITCH WHILE DO FOR GOTO CONTINUE BREAK RETURN

%start translation_unit
%%

primary_expression
        : IDENTIFIER
        | CONSTANT
        | STRING_LITERAL
        | '(' expression ')'
        ;

postfix_expression
        : primary_expression
        | postfix_expression '[' expression ']'
        | postfix_expression '(' ')'
        | postfix_expression '(' argument_expression_list ')'
        | postfix_expression '.' IDENTIFIER
        | postfix_expression PTR_OP IDENTIFIER
        | postfix_expression INC_OP
        | postfix_expression DEC_OP
        ;

argument_expression_list
        : assignment_expression
        | argument_expression_list ',' assignment_expression
        ;

unary_expression
        : postfix_expression
        | INC_OP unary_expression
        | DEC_OP unary_expression
        | unary_operator cast_expression
        | SIZEOF unary_expression
        | SIZEOF '(' type_name ')'
        ;

unary_operator
        : '&'
        | '*'
        | '+'
        | '-'
        | '~'
        | '!'
        ;

cast_expression
        : unary_expression
        | '(' type_name ')' cast_expression
        ;
```

An important question

How about C++?

Well the CFG of C++ is not LR(1).
This means we need a more advanced type of parser.

An important question

How about C++?

Well the CFG of C++ is not LR(1).
This means we need a more advanced type of parser.

(But it is <u>out-of-scope</u>).

How about Python then?

Python is SLR(1), which is – in fact – an even better class than LR(1).
This means you can use an even simpler parser than LR(1), called SLR(1) to do the job! (The acronym SLR stands for Simple LR).

(But it is <u>out-of-scope</u>).

Conclusion

- Top-Down parsing was inconclusive but gave us good insight.
- The C programming language is LR(1), we can use a bottom-up parser for the compiler of that language.
- Many other languages might not be LR(1) and might require more advanced classes of parsers, which are <u>out-of-scope</u>.
- The implementation of said parser, while interesting, is typically out-of-scope.

Conclusion

So, here is where we are.

- If our source code passes the lexical analysis and tokenizes with no trouble,

- And if there exist a derivation for the stream of tokens, that fits the YACC CFG,

- Then the code is also valid in terms of lexica and syntax (big result!).

This is not over yet, as there are still some errors that could technically make the source code malfunction, but 90% of the errors have been covered so far!

Next, semantic analysis!

Conclusion

- If interested to learn more about parsers, the reference course is the one from Stanford

https://web.stanford.edu/class/cs143/

- Available for free online and comes with video recordings

https://www.edx.org/course/compilers